Final Project: Word Embeddings and NLP in Materials Science

Aden Weiser

How can we use knowledge from natural language processing to supplement materials discovery?

LETTER

https://doi.org/10.1038/s41586-019-1335-8

Unsupervised word embeddings capture latent knowledge from materials science literature

Vahe Tshitoyan[1,3]*, John Dagdelen[1,2], Leigh Weston[1], Alexander Dunn[1,2], Ziqin Rong[1], Olga Kononova[2], Kristin A. Persson[1,2], Gerbrand Ceder[1,2]* & Anubhav Jain[1]*

Deliverable: Homework-style Jupyter Notebook

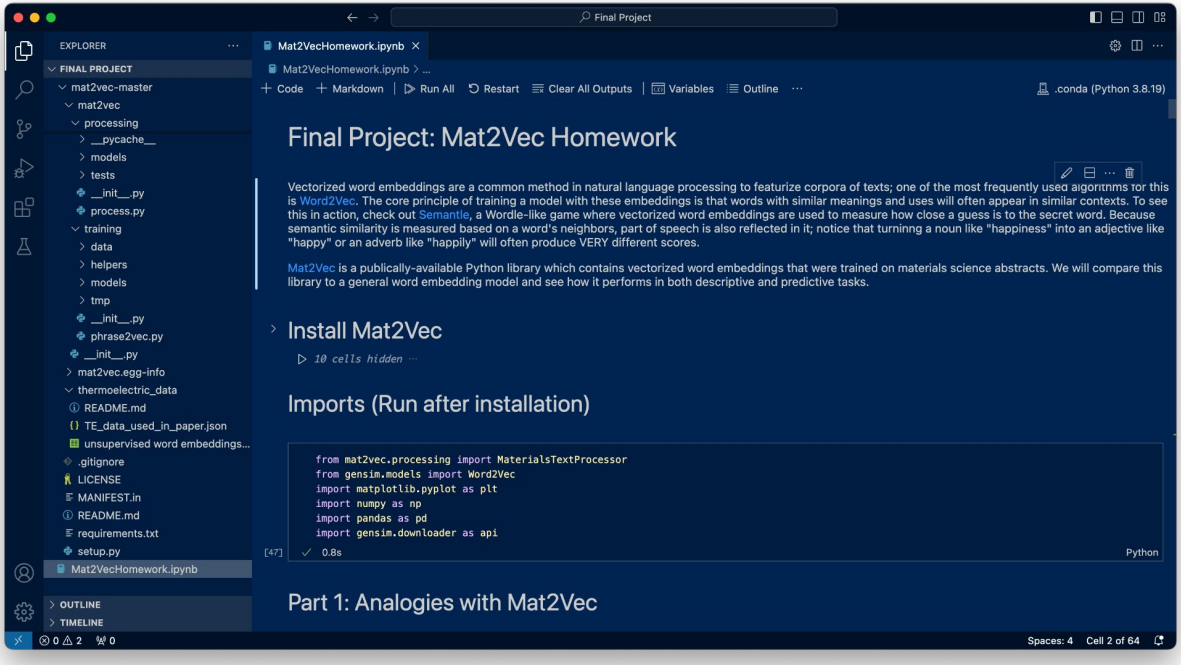

Key Learning Outcomes

- Understand basic terminology about natural language processing (corpora, tokenization, embeddings)
- Learn how training models with different data sets results in different performance on the same task
- Explore the extraction of “latent knowledge” in scientific abstracts

Task 1: Analogies with Mat2Vec only

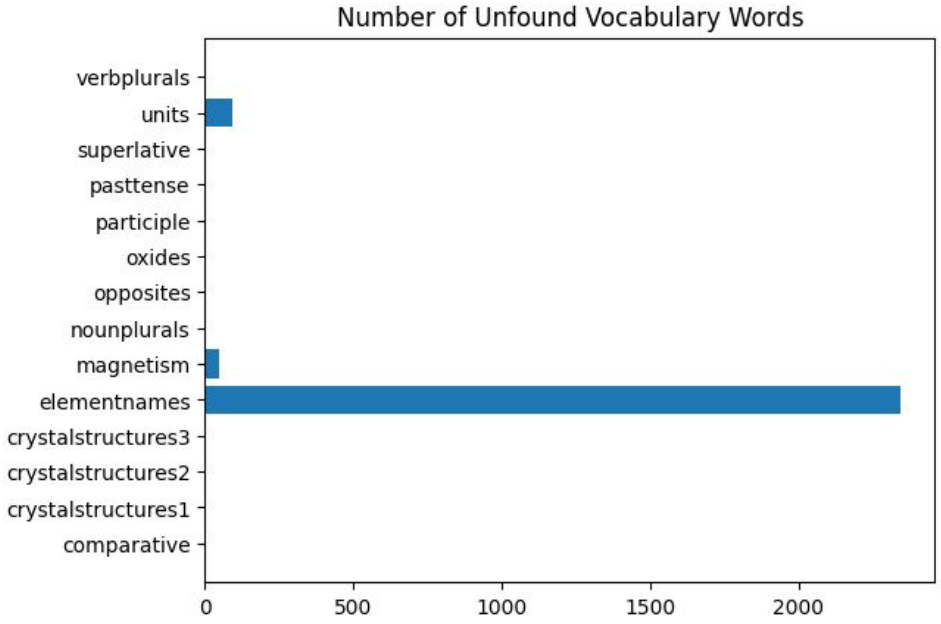

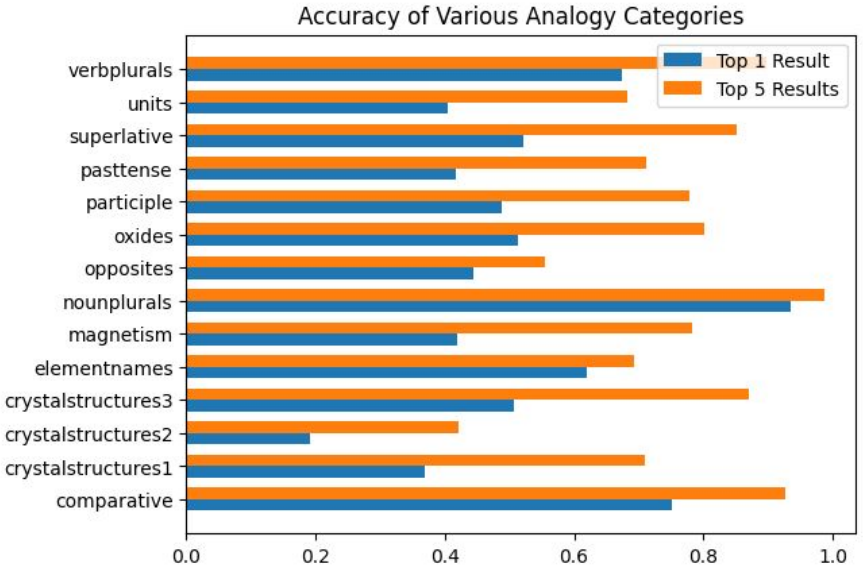

Task 2: Materials Prediction with Mat2Vec

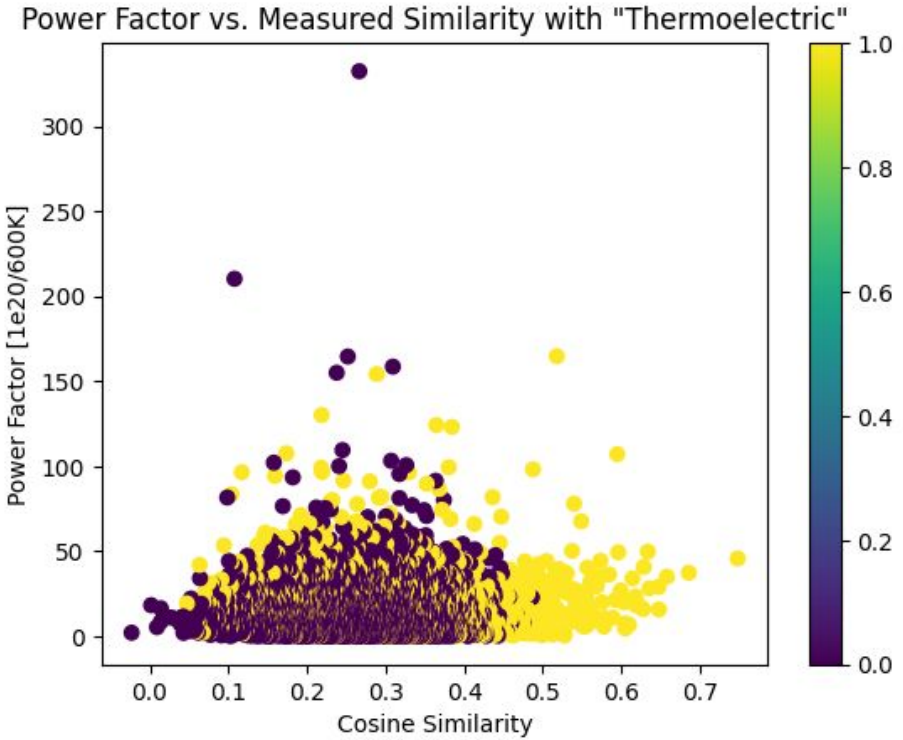

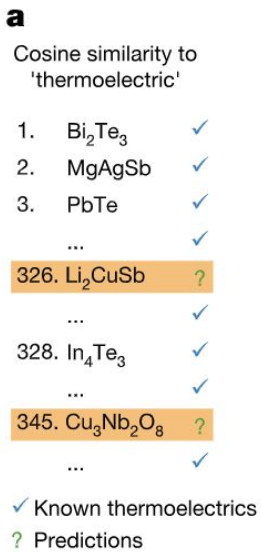

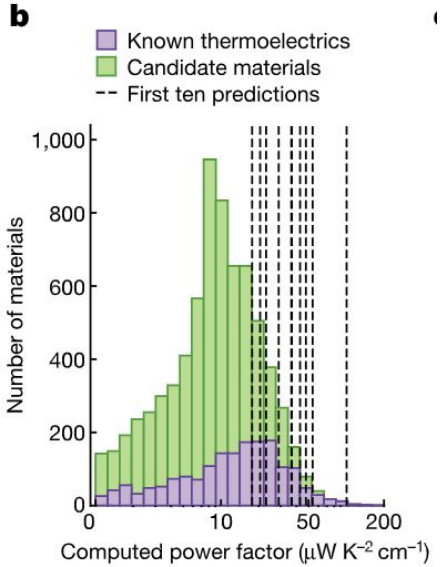

Tshitoyan et al.

Task 3: Comparison with Google News 300 Embeddings

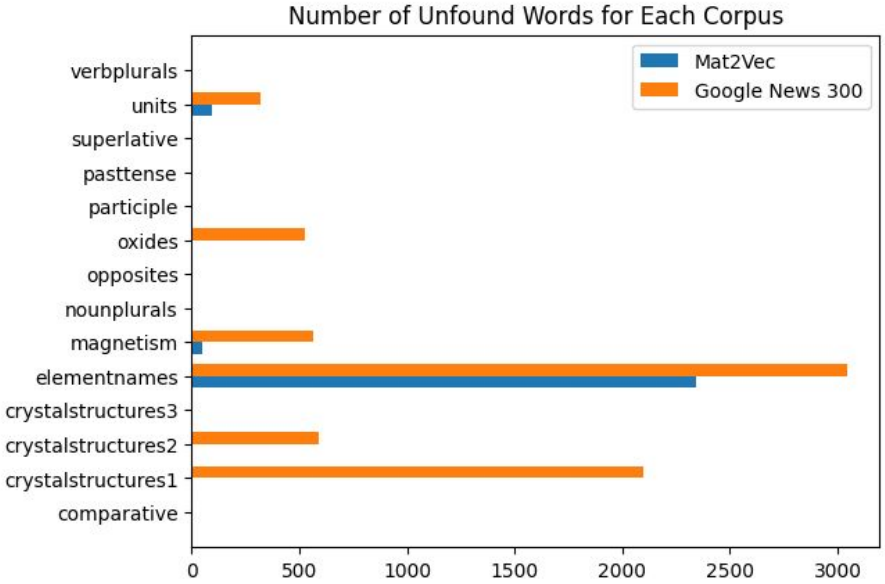

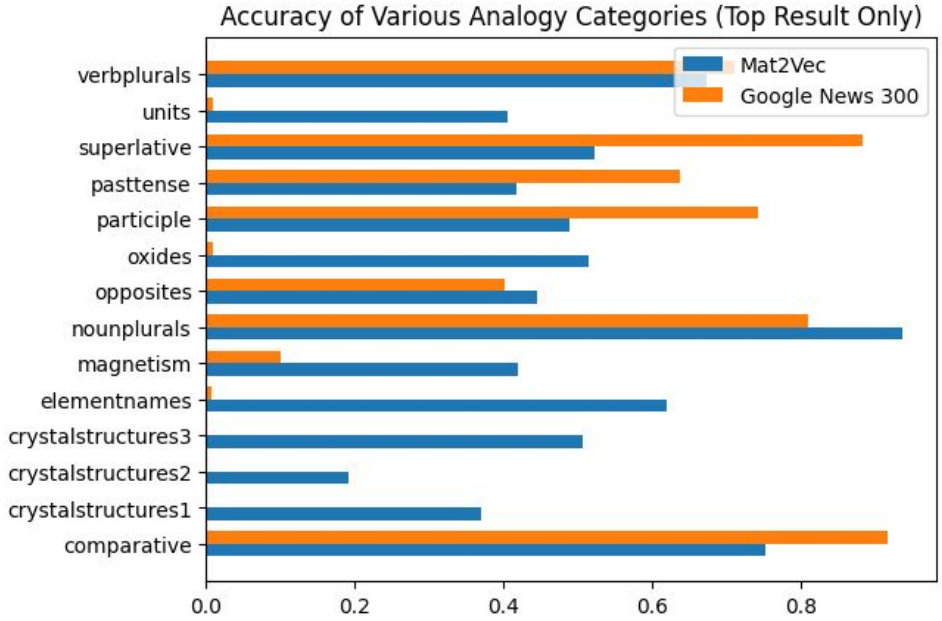

Conclusions

- Training a neural network with domain knowledge results in high performance within a domain, training a model with general knowledge results in better grammatical intuition
- Linguistic similarity is another way to featurize materials data to extract meaningful information
- Annotate your data!!!

Future Work

- Expanding/refining materials prediction task
- Integration of multiple datasets
- Google Colab integration